

Minutes of Meeting

Meeting Details

Date: February 5, 2026

Time: 4:00 PM – 4:30 PM

Platform: Google Meet

Attendees

- Client Team: Lakshman Rao Chirumamilla, Madhu Nyalam, Manish Ranjan
- Student Team: Vikesh, Advait, Amish, Keshav, Venya

Meeting Agenda

Review of technical architecture (RAG vs. Context Injection), token optimization strategies, and prototype status update.

1. Logistics and Communication

- Pre-reading Protocol: The client requested that all relevant documents, research papers, or discussion points be shared at least one day in advance (by 10:00 PM the previous night) to enable sufficient review time and more efficient discussions.

2. Technical Architecture & Research

- RAG vs. Context Injection:
 - The student team proposed a Retrieval-Augmented Generation (RAG) approach for handling medical records.
 - Decision: The client advised against implementing RAG at this stage.
 - Rationale: As the dataset consists of structured per-patient JSON files rather than large unstructured repositories (e.g., hundreds of PDFs), direct context injection was deemed sufficient and less complex for the current scope.
- Token Optimization Strategies:
 - The team raised concerns regarding inefficiencies of passing entire medical records into the LLM context window.
 - Graph-Based Inference: A research approach leveraging the National Institutes of Health (NIH) library was presented, where diagnoses are mapped to standardized codes and organized into a graph structure to enable more efficient inference. The client responded positively to this direction.

3. Project Status & Deliverables

- Prototype Update: Development was temporarily paused during the previous week due to academic evaluations (quizzes).
- Application Layer Requirement: The client emphasized that the solution must include a user-

facing application layer. The chatbot should provide an interactive chat interface for doctors or patients, and not exist solely as a backend script.

4. Action Items

Action Item	Owner	Description	Deadline
Review Research Materials	Client	Review shared research paper screenshots and the proposed NIH-based graph architecture.	Asynchronous
Finalize Baseline Prototype	Student Team	Complete a functioning prototype that ingests patient JSON and answers queries strictly from that context.	Next Meeting
Architecture Flowchart	Student Team	Formalize the shift from RAG to Context Injection in a system flowchart (completed post-meeting).	Immediate

5. Next Meeting Goals

Demonstrate a working "Baseline" chatbot through a live demo. The model should successfully answer health-related queries using only the patient-specific context provided in the JSON data.