

Minutes of Meeting

Meeting Details

Date: April 19, 2026 (Sunday)

Time: 6:00 PM

Location: CIE, IIIT-H (Offline)

Attendees

Client Team:

- Manish Ranjan (CEO, NanoCare Health Services)

Student Team:

- Vikesh
- Advait
- Amish
- Keshav
- Venya

Meeting Agenda

- Detailed walkthrough of the system architecture and RAG pipeline.
- End-to-end demonstration of all chatbot features.
- Performance analysis of local LLMs (MedGemma) vs. Gemini API.
- Privacy, GDPR compliance, and PII de-identification implementation.
- Planning for beta testing and external feedback gathering.
- Discussion on the industry trade-off between Accuracy and Optimization.

1. System Walkthrough and Demonstration

- The student team provided a comprehensive explanation of the entire RAG pipeline, guiding the client through the previously shared architectural presentation.
- A detailed, feature-by-feature live demonstration of the Robert AI Medical Assistant was conducted, showcasing EMR ingestion, intent classification, and clinical response generation.

2. Local Model Performance and Hardware Optimization

- The system's inference capabilities were demonstrated using both the cloud-based Gemini API and the local MedGemma model.
- The client noted that MedGemma experienced significant load times and latency.
- The team discussed whether the latency bottleneck was a compute-bound or memory-bound issue.
- Hardware acceleration possibilities were explored, specifically the potential of running the model optimally on dedicated GPUs or an Apple Mac Mini to improve throughput.

3. Privacy, Testing, and Real-World Deployment

- The team reviewed the implemented PII de-identification mechanisms and GDPR compliance measures.
- To ensure the project's practical utility, the client proposed an outsourced testing phase. The system will be deployed to his internal team and an additional small group of external users to gather critical, real-world feedback.

4. Core Industry Challenge: Accuracy vs. Optimization

- The client asked for the team's perspective regarding a central problem in the healthcare software industry: the trade-off between maximizing diagnostic accuracy and maintaining computational optimization when processing massive Electronic Medical Records (EMRs).

5. Action Items

Action Item	Owner	Description	Status/Deadline
Implement MPS GPU Backend Support	Student Team	Add support for Apple Metal Performance Shaders (MPS) as a GPU backend to optimize local model execution.	Completed (Post-Meet)
Beta Deployment Setup	Client & Team	Prepare the system infrastructure to allow the client's internal team and external testing group to use and review the chatbot.	Next Phase